

Customer Churn Prediction

July 19, 2025

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn import preprocessing
```

```
/Users/hinalpanchal/opt/anaconda3/lib/python3.9/site-
packages/scipy/_init_.py:155: UserWarning: A NumPy version >=1.18.5 and
<1.25.0 is required for this version of SciPy (detected version 1.26.4
  warnings.warn(f"A NumPy version >={np_minversion} and <{np_maxversion}")
```

```
[2]: telecom_data = pd.read_csv('Customer_Churn.csv')
telecom_data.head()
```

```
[2]:
```

	CustomerID	Age	Gender	Tenure	MonthlyCharges	ContractType	\
0	1	49	Male	4	88.35	Month-to-Month	
1	2	43	Male	0	36.67	Month-to-Month	
2	3	51	Female	2	63.79	Month-to-Month	
3	4	60	Female	8	102.34	One-Year	
4	5	42	Male	32	69.01	Month-to-Month	

	InternetService	TotalCharges	TechSupport	Churn
0	Fiber Optic	353.40	Yes	Yes
1	Fiber Optic	0.00	Yes	Yes
2	Fiber Optic	127.58	No	Yes
3	DSL	818.72	Yes	Yes
4	None	2208.32	No	Yes

```
[3]: telecom_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 10 columns):
#   Column          Non-Null Count  Dtype
---  -
0   CustomerID      1000 non-null   int64
1   Age             1000 non-null   int64
2   Gender          1000 non-null   object
3   Tenure          1000 non-null   int64
```

```
4   MonthlyCharges    1000 non-null    float64
5   ContractType      1000 non-null    object
6   InternetService   1000 non-null    object
7   TotalCharges      1000 non-null    float64
8   TechSupport       1000 non-null    object
9   Churn             1000 non-null    object
dtypes: float64(2), int64(3), object(5)
memory usage: 78.2+ KB
```

```
[4]: telecom_data.isna().sum()
```

```
[4]: CustomerID      0
Age                0
Gender             0
Tenure            0
MonthlyCharges    0
ContractType      0
InternetService   0
TotalCharges      0
TechSupport       0
Churn             0
dtype: int64
```

```
[5]: telecom_data.duplicated().sum()
```

```
[5]: 0
```

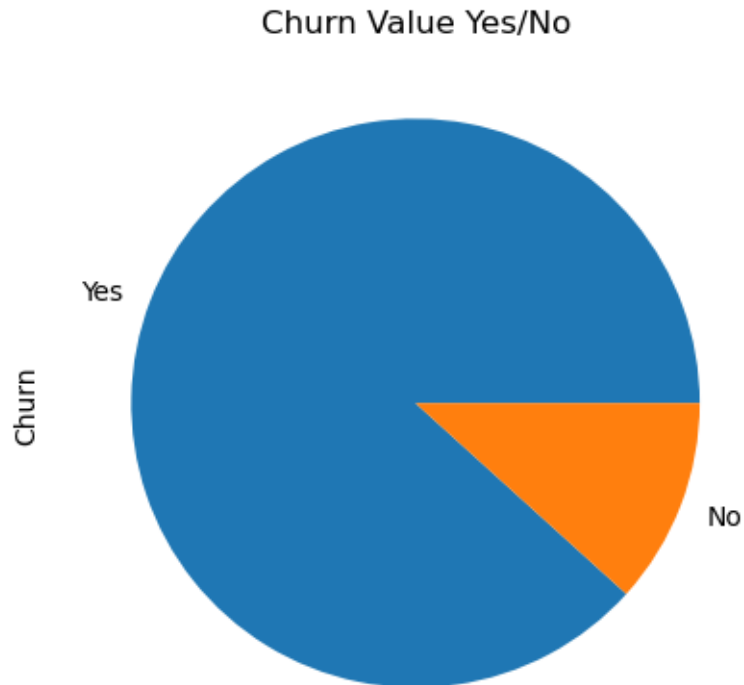
1 EDA

```
[6]: telecom_data['Churn'].value_counts()
```

```
[6]: Yes      883
No       117
Name: Churn, dtype: int64
```

```
[7]: telecom_data['Churn'].value_counts().plot(kind='pie')
plt.title('Churn Value Yes/No')
```

```
[7]: Text(0.5, 1.0, 'Churn Value Yes/No')
```



```
[8]: telecom_data.groupby('Churn')['MonthlyCharges'].mean()
```

```
[8]: Churn
     No      62.54641
     Yes      75.96077
     Name: MonthlyCharges, dtype: float64
```

```
[9]: telecom_data.groupby(['Churn', 'Gender'])['MonthlyCharges'].mean()
```


```
[9]: Churn  Gender
     No   Female    65.091912
         Male     59.013878
     Yes   Female    74.975064
         Male     77.082518
     Name: MonthlyCharges, dtype: float64
```

```
[10]: telecom_data.groupby('Churn')['Tenure'].mean()
```

```
[10]: Churn
     No      30.264957
     Yes     17.476784
     Name: Tenure, dtype: float64
```

```
[11]: telecom_data.groupby('ContractType')['MonthlyCharges'].mean().plot(kind='bar')
plt.ylabel('Mean Price')
plt.xlabel('Contract Type')
```

```
[11]: Text(0.5, 0, 'Contract Type')
```




```
[12]: telecom_data.groupby(['InternetService', 'Churn'])['TotalCharges'].mean()
```

```
[12]: InternetService  Churn
DSL                  No      1650.001875
                   Yes      1342.233038
Fiber Optic         No      1976.429130
                   Yes      1288.080798
None                 Yes      1413.789327
Name: TotalCharges, dtype: float64
```


```
[13]: plt.hist(telecom_data['MonthlyCharges'])  
plt.title('Monthly Charges')
```

```
[13]: Text(0.5, 1.0, 'Monthly Charges')
```



```
[14]: plt.hist(telecom_data['Tenure'])  
plt.title('Histogram for Tenure')
```

```
[14]: Text(0.5, 1.0, 'Histogram for Tenure')
```



2 Feature Re-Engineering

```
[15]: from sklearn.model_selection import train_test_split
      from sklearn.preprocessing import StandardScaler
      from sklearn.linear_model import LogisticRegression
      from sklearn.metrics import classification_report, confusion_matrix
      from sklearn.pipeline import Pipeline
```

```
[16]: telecom_data.head(1)
```

```
[16]:   CustomerID  Age  Gender  Tenure  MonthlyCharges  ContractType \
0           1   49   Male      4           88.35  Month-to-Month

   InternetService  TotalCharges  TechSupport  Churn
0      Fiber Optic          353.4         Yes    Yes
```

```
[17]: telecom_data.columns
```

```
[17]: Index(['CustomerID', 'Age', 'Gender', 'Tenure', 'MonthlyCharges',
        'ContractType', 'InternetService', 'TotalCharges', 'TechSupport',
```

```
    'Churn'],  
    dtype='object')
```

```
[18]: y= telecom_data['Churn']  
      X= telecom_data[['Age', 'Gender', 'Tenure', 'MonthlyCharges']]
```

```
[19]: # Gender 1- Female, 0-Male  
      X.loc[:, 'Gender'] = X['Gender'].apply(lambda x: 1 if x == 'Female' else 0)  
  
/var/folders/qw/sl7j3jj94ngg0_334hjh2ycm0000gn/T/ipykernel_44305/1439695941.py:2  
: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead  
  
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy  
      X.loc[:, 'Gender'] = X['Gender'].apply(lambda x: 1 if x == 'Female' else 0)
```

```
[20]: X.head(3)
```

```
[20]:
```

	Age	Gender	Tenure	MonthlyCharges
0	49	0	4	88.35
1	43	0	0	36.67
2	51	1	2	63.79

```
[21]: # Churn Yes=1 ,No=0  
      # Using apply + lambda  
      y = y.apply(lambda x: 1 if x == 'Yes' else 0)
```

```
[22]: y.head(10)
```

```
[22]:
```

0	1
1	1
2	1
3	1
4	1
5	1
6	1
7	1
8	0
9	1

Name: Churn, dtype: int64

```
[23]: X_train, X_test, y_train, y_test = train_test_split(X,y, test_size=0.2)
```

```
[24]: scaler = StandardScaler()  
  
      X_train =scaler.fit_transform(X_train)
```

```
# We dump the data as before Scaling the X_test as there might be leakage when  
↪ performing Standard Scaler
```

```
[25]: import joblib  
      joblib.dump(scaler, 'scaler.pkl')
```

```
[25]: ['scaler.pkl']
```

```
[26]: X_test = scaler.fit_transform(X_test)
```

```
[32]: from sklearn.metrics import accuracy_score  
  
      def modelperformance(prediction):  
          print('Accuracy score on model is {}'.format(accuracy_score(y_test, y_pred)))
```

```
[33]: from sklearn.model_selection import GridSearchCV  
      import warnings  
      warnings.filterwarnings('ignore')
```

3 LOGISTIC REGRESSION

```
[34]: LogR = LogisticRegression()  
      LogR.fit(X_train, y_train)
```

```
[34]: LogisticRegression()
```

```
[35]: y_pred = LogR.predict(X_test)
```

```
[36]: modelperformance(y_pred)
```

Accuracy score on model is 0.88

4 KNN

```
[37]: from sklearn.neighbors import KNeighborsClassifier
```

```
[44]: param_grid = {  
      'n_neighbors': [3, 5, 7, 9],  
      'weights': ['uniform', 'distance'],  
      }  
  
      gridkn = GridSearchCV(KNeighborsClassifier(), param_grid, cv=5)
```

```
[45]: gridkn.fit(X_train, y_train)
```



```
[45]: GridSearchCV(cv=5, estimator=KNeighborsClassifier(),
                param_grid={'n_neighbors': [3, 5, 7, 9],
                            'weights': ['uniform', 'distance']})
```

```
[46]: gridkn.best_params_
```

```
[46]: {'n_neighbors': 9, 'weights': 'uniform'}
```

```
[47]: y_preds = gridkn.predict(X_test)
```

```
[48]: modelperformance(y_pred)
```

Accuracy score on model is 0.88

5 SVM

```
[49]: from sklearn.svm import SVC
```

```
[50]: svm =SVC()
      param_grid_svm = {'C':[0.01,0.1,0.5,1], 'kernel':['linear','rbf','poly']}
```

```
[51]: gridsvm = GridSearchCV(svm,param_grid_svm,cv=5)
```

```
[52]: gridsvm.fit(X_train,y_train)
```

```
[52]: GridSearchCV(cv=5, estimator=SVC(),
                param_grid={'C': [0.01, 0.1, 0.5, 1],
                            'kernel': ['linear', 'rbf', 'poly']})
```

```
[53]: gridsvm.best_params_
```

```
[53]: {'C': 0.01, 'kernel': 'linear'}
```

```
[54]: y_pred = gridsvm.predict(X_test)
```

```
[55]: modelperformance(y_pred)
```

Accuracy score on model is 0.89

6 Decision Tree Classifier

```
[56]: from sklearn.tree import DecisionTreeClassifier

      param_grid ={
          'criterion': ["entropy", "gini"],
          'splitter': ['best','random'],
          'max_depth': [None,10,20,30],
          'min_samples_leaf': [1, 2, 4],
```

```
    'min_samples_split': [2, 5, 10]
}
```

```
[57]: grid_tree = GridSearchCV(DecisionTreeClassifier(),param_grid, cv=5)
```

```
[58]: grid_tree.fit(X_train,y_train)
```

```
[58]: GridSearchCV(cv=5, estimator=DecisionTreeClassifier(),
                  param_grid={'criterion': ['entropy', 'gini'],
                              'max_depth': [None, 10, 20, 30],
                              'min_samples_leaf': [1, 2, 4],
                              'min_samples_split': [2, 5, 10],
                              'splitter': ['best', 'random']})
```

```
[59]: grid_tree.best_params_
```

```
[59]: {'criterion': 'entropy',
      'max_depth': None,
      'min_samples_leaf': 4,
      'min_samples_split': 10,
      'splitter': 'random'}
```

```
[60]: y_pred = grid_tree.predict(X_test)
```

```
[61]: modelperformance(y_pred)
```

Accuracy score on model is 0.805

7 Random Forest Classifier

```
[63]: from sklearn.ensemble import RandomForestClassifier
```

```
[65]: rfc_model = RandomForestClassifier()

param_grid = {
    'n_estimators': [32,64,128,256],
    'max_features': [2,3,4],
    'bootstrap': [True,False]
}
```

```
[66]: rfc = GridSearchCV(rfc_model,param_grid,cv=5)
```

```
[67]: rfc.fit(X_train,y_train)
```

```
[67]: GridSearchCV(cv=5, estimator=RandomForestClassifier(),
                  param_grid={'bootstrap': [True, False], 'max_features': [2, 3, 4],
                              'n_estimators': [32, 64, 128, 256]})
```

```
[70]: rfc.best_params_
```

```
[70]: {'bootstrap': True, 'max_features': 2, 'n_estimators': 128}
```

```
[68]: y_pred = rfc.predict(X_test)
```

```
[69]: modelperformance(y_pred)
```

Accuracy score on model is 0.885

```
[71]: best_model = gridsvm.best_estimator_
```

```
[72]: joblib.dump(best_model, 'model.pkl')
```

```
[72]: ['model.pkl']
```

```
[ ]:
```